

SCALABLE AND RELIABLE VIDEO TRANSCODING IN HASKELL

Alfredo Di Napoli

Haskell Exchange 2015

Full story at: <http://goo.gl/qkKwKm>



```
launchMissile :: IO ()
```



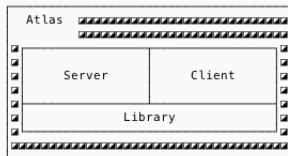


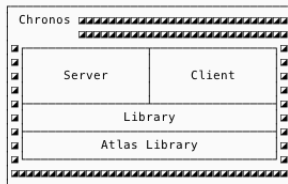
iris
Change. Doing it right.



Market leader for CPD solutions

Used by more than 1000 schools across UK, Europe, US & Australia

IRIS' GREEK ZOO



October, 2013

```
hermes [74aeab2] ⚡ cloc src scripts
  3 text files.
  3 unique files.
  0 files ignored.

http://cloc.sourceforge.net v 1.60 T=0.24 s (12.3 files/s, 762.9 lines/s)
-----
Language          files          blank          comment          code
-----
Haskell              1              22              7              80
Bourne Shell         2              25              2              50
-----
SUM:                  3              47              9             130
-----
```

October, 2015

```
hermes [master] cloc main server src test
 105 text files.
 105 unique files.
  0 files ignored.

http://cloc.sourceforge.net v 1.60 T=0.66 s (159.6 files/s, 26288.3 lines/s)
-----
Language          files          blank          comment          code
-----
Haskell            105           2332           2252           12712
-----
SUM:                105           2332           2252           12712
-----
```

Upon taking the lead on Hermes, I was asked for a couple of requirements to be fulfilled, the most important one being that the system needed to be deployed in a cluster, capable of scaling according to demand.

More specifically, we wanted a system with these desirable properties:

- ~ Scalable
- ~ Fault tolerant
- ~ Distributed

The classical ways to approach the difficulty of state include OOP programming which tightly couples state together with related behaviour, and functional programming which — in its pure form — eschews state and side-effects all together. [..] We argue that it is possible to take useful ideas from both and that this approach offers significant potential for simplifying the construction of large-scale software systems.

In the same fashion, we have 2 different worlds colliding:

- ~ We need to transcode videos, which is a very stateful operation
- ~ As good Haskell programmers, we want to have components in our system to be as stateless as possible.

A shared nothing architecture (SN) is a distributed computing architecture in which each node is independent and self-sufficient, and there is no single point of contention across the system.

"All problems in computer science can be solved by another level of indirection." - Butler Lampson

1. RabbitMQ was just the right tool for the job at hand:
 - ~ Easy to setup
 - ~ Can be configured to operate in a federation of nodes
 - ~ Extremely reliable
 - ~ Good Haskell bindings for it (*AMQP*)
2. A question genuinely arise: it seems extremely costly to shuffle video as binary blobs over the queues. Can we avoid that?

ABSTRACTION IS THE (MEDIA KEY)

```
root__m-stg-main-2014_10_29_13_27_26-videos-1-2333-vid-smc-oxz8dmdillx7fong
  ^      ^      ^      ^      ^      ^      ^      ^      ^      ^
  |      |      |      |      |      |      |      |      |      |
comment  |      |      |      |      |      |      |      |      |
host ----+      |      |      |      |      |      |      |      |
database --- +   |      |      |      |      |      |      |      |
dataset version -----+  |      |      |      |      |      |      |
resource (video or image) -----+  |      |      |      |      |      |
user ID  -----+  |      |      |      |      |      |      |      |
video ID -----+  |      |      |      |      |      |      |      |
channel type -----+  |      |      |      |      |      |      |      |
video products -----+  |      |      |      |      |      |      |      |
MAC (avoids submission of bogus keys) -----+  |      |      |      |      |      |      |
```


To be fair, the media key abstraction was already present in Atlas when I choose RabbitMQ, but it was the perfect fit for it!

Fine, but RabbitMQ doesn't give you scalability...

1. We stood on the shoulder of giants - namely AWS' Auto Scaling Groups
2. Our very first native scaling algorithm looked like:
 - ~ Scaling up: Based on CPU% over time
 - ~ Scaling down: Based on CPU% over time

It kept us going for a while...

THE ARCHITECTURE



1. Scaling up was too conservative and slow
 - ~ It could take up to 15 mins to spawn a new worker
2. Scaling down suffered similar problems
3. The result was unoptimal customer experience (due to the slow turnaround time) and unoptimal for us (due to the additional costs incurring from poor scaling down)

Overview EC2 Cost Throughput Efficiency CPU Utilisation Cluster Size

These charts represent the cost of every EC2 instance labeled as an "Hermes machine", namely hermes-stg/hermes-dev, the regionalised hermes servers and the regionalised workers.

EU Region Production

☐ Server ☐ Worker ☒ All ☐ Invalidate cache

From:

2015-06-01

Until:


2015-06-30

Chart 1

EC2 Cost

● c1.medium ● c3.2xlarge ● c3.large ● c3.xlarge



EC2 Cost

● c1.medium ● c3.2xlarge ● c3.large ● c3.xlarge



EC2 Cost





THE ELEPHANT IN THE ROOM

What's the elephant in the room?



Why not use Cloud Haskell?

1. CH encourages Erlang-style (i.e. actor based) communication, so nodes should know each other

We do not want that!

2. Peer discovery would have been tricky in a dynamic environment where new machines born and die frequently
3. It wasn't mature enough in 2013, if not for a handful of companies using it

Thank you!

Questions?

My road to Haskell

<http://www.alfredodinapoli.com/posts/2014-04-27-my-road-to-haskell.html>

Don Stewart - Haskell in the large

<http://code.haskell.org/~dons/talks/dons-google-2015-01-27.pdf>