

SCALABLE AND RELIABLE VIDEO TRANSCODING IN HASKELL

Alfredo Di Napoli

Haskell Exchange 2015

Full story at: <http://goo.gl/qkKwKm>



```
launchMissile :: IO ()
```








discover. develop. share.

- ~ Present in over 1800 schools Worldwide (mostly UK, Europe, US & Australia)
- ~ Used by over 32000 teachers

IRIS CONNECT (CONTD.)

 Home Reflections Forms Groups


Search  

New Reflection +

My Reflections >



Shared with me >


Invitations >



Mandy's Lesson Starters
Demo User 02/15/15 11:14



Review


Details  



History Lesson
Demo User 03/15/15 12:27



Review


Details  



Dual View
Demo User 01/20/15 12:34



Review


Details  



Cramlington Learning Village
Demo User 09/22/14 12:25



Review


Details  



Maths Lesson
Demo User 09/09/14 14:23



Review


Details  



History Lesson Cartoon
Demo User 06/20/15 12:34



Review


Details  



Questioning Best Practice
Demo User 03/21/15 12:25



Review

Details  



LiveView Sample - Math Lesson
Demo User 04/10/15 12:43

Review

Details  

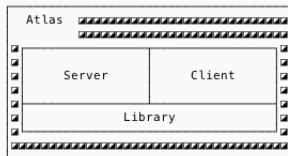


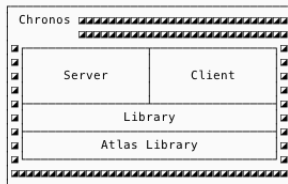
LiveView



Discovery Kit



IRIS' GREEK ZOO



HERMES' EVOLUTION

October, 2013

```
hermes [74aeab2] ⚡ cloc src scripts
  3 text files.
  3 unique files.
  0 files ignored.

http://cloc.sourceforge.net v 1.60 T=0.24 s (12.3 files/s, 762.9 lines/s)
-----
Language          files          blank          comment          code
-----
Haskell              1              22              7              80
Bourne Shell         2              25              2              50
-----
SUM:                 3              47              9             130
-----
```

October, 2015

```
hermes [master] cloc main server src test
 105 text files.
 105 unique files.
  0 files ignored.

http://cloc.sourceforge.net v 1.60 T=0.66 s (159.6 files/s, 26288.3 lines/s)
-----
Language          files          blank          comment          code
-----
Haskell           105          2332          2252          12712
-----
SUM:              105          2332          2252          12712
-----
```

Upon taking the lead on Hermes, I was asked for a couple of requirements to be fulfilled, the most important one being that the system needed to be deployed in a cluster, capable of scaling according to demand.

More specifically, we wanted a system with these desirable properties:

- ~ Scalable
- ~ Fault tolerant
- ~ Distributed

- ~ It's easy to see that what we want is a **cluster**, capable of scaling on demand
- ~ We need to transcode videos, which is a very stateful operation
- ~ A cluster typically implies machines talking to each other, which is also very stateful
- ~ **As good Haskell programmers**, we want to have components in our system to be **as stateless as possible**, and potentially treat videos as **persistent data structures**!

A shared nothing architecture (SN) is a distributed computing architecture in which each node is independent and self-sufficient, and there is no single point of contention across the system.

"All problems in computer science can be solved by another level of indirection." - Butler Lampson

1. RabbitMQ was just the right tool for the job at hand:
 - ~ Easy to setup
 - ~ Can be configured to operate in a federation of nodes
 - ~ Extremely reliable
 - ~ Good Haskell bindings for it (*AMQP*)
2. A question genuinely arise: it seems extremely costly to shuffle video as binary blobs over the queues. Can we avoid that?

ABSTRACTION IS THE (MEDIA KEY)

```
root__m-stg-main-2014_10_29_13_27_26-videos-1-2333-vid-smc-oxz8dmdillx7fong
  ^      ^      ^      ^      ^      ^      ^      ^      ^      ^
  |      |      |      |      |      |      |      |      |
comment  |      |      |      |      |      |      |      |
host ----+      |      |      |      |      |      |      |
database --- +   |      |      |      |      |      |      |
dataset version -----+      |      |      |      |      |
resource (video or image) -----+      |      |      |      |
user ID -----+      |      |      |      |      |      |
video ID -----+      |      |      |      |      |      |
channel type -----+      |      |      |      |      |
video products -----+      |      |      |      |      |
MAC (avoids submission of bogus keys) -----+      |      |
```

To be fair, the media key abstraction was already present in Atlas when I choose RabbitMQ, but it was the perfect fit for it!

WHAT ABOUT DATA STORAGE?

Fine, but RabbitMQ doesn't give you data persistence...

1. We use AWS' S3 for our storing needs

~ A media key **uniquely identifies** an S3 location (it's like an **IP address for videos!**)

2. Upon upload the original file from the user is synced over S3 and we call this generation-0 file the **master file**
3. Such master file is **immutable**, and each product we transcode generates a brand new binary on S3

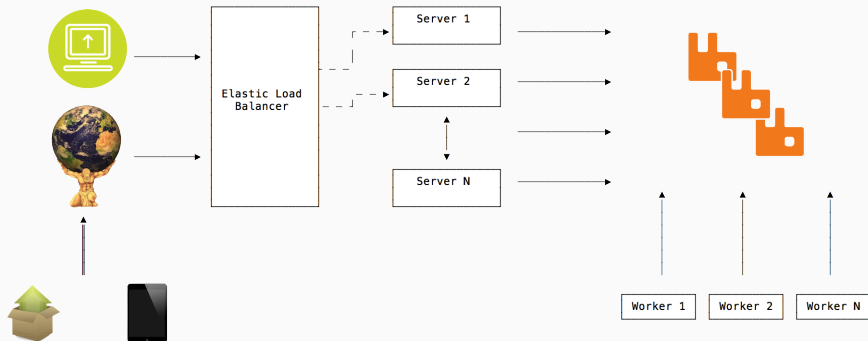
We are treating videos as immutable data structures!

Fine, but RabbitMQ doesn't give you scalability...

1. We stood once again on the shoulder of giants - namely AWS' Auto Scaling Groups
2. Our very first native scaling algorithm looked like:
 - ~ Scaling up: Based on CPU% over time
 - ~ Scaling down: Based on CPU% over time

It kept us going for a while...

THE ARCHITECTURE



1. Scaling up was too conservative and slow
 - ~ It could take up to 15 mins to spawn a new worker
2. Scaling down suffered similar problems
3. The result was unoptimal customer experience (due to the slow turnaround time) and unoptimal for us (due to the additional costs incurring from poor scaling down)

Overview EC2 Cost Throughput Efficiency CPU Utilisation Cluster Size

These charts represent the cost of every EC2 Instance labeled as an "Hermes machine", namely hermes-stg/hermes-dev, the regionalised hermes servers and the regionalised workers.

EU Region

Production

☐ Server ☐ Worker ☒ All ☐ Invalidate cache

From:

2015-06-01

Until:


2015-06-30

Chart 11

EC2 Cost

● c1.medium ● c3.2xlarge ● c3.large ● c3.xlarge



EC2 Cost

● c1.medium ● c3.2xlarge ● c3.large ● c3.xlarge



EC2 Cost





THE ELEPHANT IN THE ROOM

What's the elephant in the room?



Why not use Cloud Haskell?

1. CH encourages Erlang-style (i.e. actor based) communication, so nodes should know each other

We do not want that!

2. Peer discovery would have been tricky in a dynamic environment where new machines born and die frequently
3. It wasn't mature enough in 2013, if not for a handful of companies using it

Thank you!

Questions?

My road to Haskell

<http://www.alfredodinapoli.com/posts/2014-04-27-my-road-to-haskell.html>

Don Stewart - Haskell in the large

<http://code.haskell.org/~dons/talks/dons-google-2015-01-27.pdf>